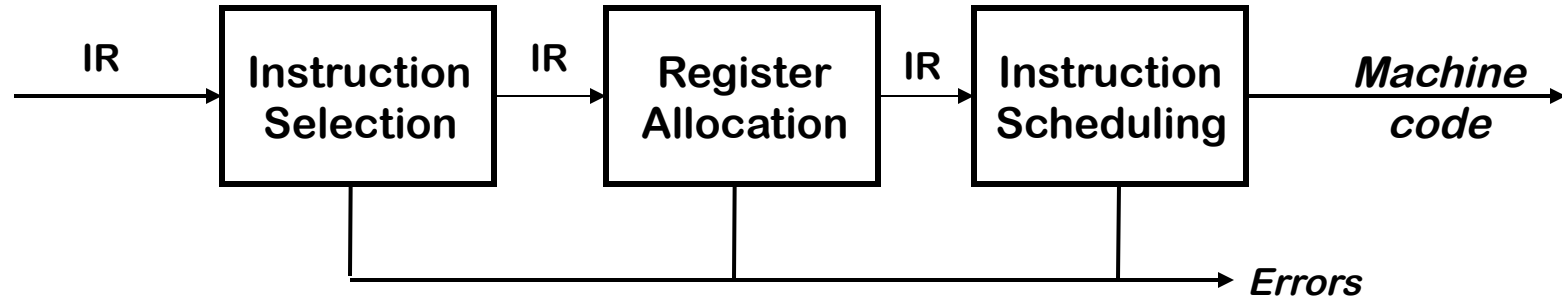


CS415 Compilers

Instruction Scheduling and Introduction to Lexical Analysis

These slides are based on slides copyrighted by
Keith Cooper, Ken Kennedy & Linda Torczon at Rice
University



Responsibilities

- Translate IR into target machine code
- Choose instructions to implement each IR operation
- Decide which value to keep in registers
- Ensure conformance with system interfaces

Automation has been *less* successful in the back end

Readings: Eac 12.1-12.3

Local: within single basic block

Global: across entire procedure

Motivation

- Instruction latency (pipelining)
several cycles to complete instructions; instructions can be issued every cycle
- Instruction-level parallelism (VLIW, superscalar)
execute multiple instructions per cycle

Issues

- Reorder instructions to reduce execution time (or power requirements)
- Static schedule - insert NOPs to preserve correctness
- Dynamic schedule - hardware pipeline stalls
- Preserve correctness
- Interactions with other optimizations (register allocation!)

Motivation

- Instruction latency (pipelining)
several cycles to complete instructions; instructions can be issued every cycle
- Instruction-level parallelism (VLIW, superscalar)
execute multiple instructions per cycle

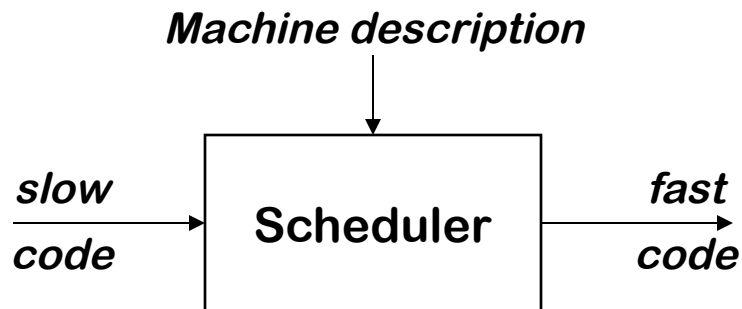
Issues

- Reorder instructions to reduce execution time (or power requirements)
- Static schedule - insert NOPs to preserve correctness
- Dynamic schedule - hardware pipeline stalls
- Preserve correctness
- Interactions with other optimizations (register allocation!)
- Note: code shape contains real, not virtual registers
==> register may be redefined

The Problem

Given a code fragment for some target machine and the latencies for each individual operation, reorder the operations to minimize execution time

The Concept



The task

- Produce correct code
- Minimize wasted (idle) cycles
- Operate efficiently

Dependences \Rightarrow defined on memory locations / registers and not values

Statement/instruction **b** depends on statement/instruction **a** if there exists:

- **true** of flow dependence
 a writes a location/register that **b** later reads (RAW conflict)
- **anti** dependence
 a reads a location/register that **b** later writes (WAR conflict)
- **output** dependence
 a writes a location/register that **b** later writes (WAW conflict)

Dependences define ORDER CONSTRAINTS that need to be respected in order to generate correct code.


true	anti	output
a =	= a	a =
= a	a =	a =

To capture properties of the code, build a dependence graph G

- Nodes $n \in G$ are operations with $type(n)$ and $delay(n)$
- An edge $e = (n_1, n_2) \in G$ iff n_2 uses the result of n_1

a:	loadAl	r0,@w	\Rightarrow r1
b:	add	r1,r1	\Rightarrow r1
c:	loadAl	r0,@x	\Rightarrow r2
d:	mult	r1,r2	\Rightarrow r1
e:	loadAl	r0,@y	\Rightarrow r3
f:	mult	r1,r3	\Rightarrow r1
g:	loadAl	r0,@z	\Rightarrow r2
h:	mult	r1,r2	\Rightarrow r1
i:	storeAl	r1	\Rightarrow r0,@w

The Code



The Dependence Graph

(all output dependences are covered,
i.e., are satisfied through other
dependences)

A correct schedule S maps each $n \in N$ into a non-negative integer representing its cycle number such that

1. $S(n) \geq 0$, for all $n \in N$, obviously
2. If $(n_1, n_2) \in E$, $S(n_1) + \text{delay}(n_1) \leq S(n_2)$
3. For each type t , there are no more operations of type t in any cycle than the target machine can issue

The length of a schedule S , denoted $L(S)$, is

$$L(S) = \max_{n \in N} (S(n) + \text{delay}(n))$$

The goal is to find the shortest possible correct schedule.

S is time-optimal if $L(S) \leq L(S_1)$, for all other schedules S_1

A schedule might also be optimal in terms of registers, power, or space....

Critical Points

- All operands must be available
- Multiple operations can be ready
- Moving operations can lengthen register lifetimes
- Placing uses near definitions can shorten register lifetimes
- Operands can have multiple predecessors

Together, these issues make scheduling hard (NP-Complete)

Local scheduling is the simpler case

- Restricted to straight-line code (single basic block)
- Consistent and predictable latencies

The big picture

1. Build a dependence graph, P
2. Compute a priority function over the nodes in P
3. Use list scheduling to construct a schedule, one cycle at a time
(can only issue/schedule at most one instructions per cycle)
 - a. Use a queue of operations that are ready
 - b. At each cycle
 - I. Choose a ready operation and schedule it
 - II. Update the ready queue

Local list scheduling

- The dominant algorithm for twenty years
- A greedy, heuristic, local technique

```
Cycle  $\leftarrow$  1
Ready  $\leftarrow$  leaves of  $P$ 
Active  $\leftarrow \emptyset$ 

while (Ready  $\cup$  Active  $\neq \emptyset$ )
  if (Ready  $\neq \emptyset$ ) then
    remove an  $op$  from Ready
     $S(op) \leftarrow$  Cycle
    Active  $\leftarrow$  Active  $\cup op$ 

  Cycle  $\leftarrow$  Cycle + 1

  for each  $op \in$  Active
    if ( $S(op) + \text{delay}(op) \leq$  Cycle) then
      remove  $op$  from Active
      for each successor  $s$  of  $op$  in  $P$ 
        if ( $s$  is ready) then
          Ready  $\leftarrow$  Ready  $\cup s$ 
```

Removal in priority order

op has completed execution

If successor's operands are ready, put it on Ready

<u>Operation</u>	<u>Cycles</u>
load	3
loadl	1
loadAl	3
store	3
storeAl	3
add	1
mult	2
fadd	1
fmult	2
shift	1
branch	0 to 8

- **Loads & stores may or may not block**
 - > Non-blocking \Rightarrow fill those issue slots
- **Branches typically have delay slots**
 - > Fill slots with operations unrelated to branch condition evaluation
 - > Percolates branch upward
- **Branch Prediction may hide branch latencies (hardware feature)**

Build a simple local scheduler (basic block)


- non-blocking loads & stores
- different latencies load/store vs. arith. etc. operations
- different heuristics
- forward / backward scheduling

1. Build the dependence graph

1	a:	loadAl	r0,@w	⇒ r1
4	b:	add	r1,r1	⇒ r1
5	c:	loadAl	r0,@x	⇒ r2
8	d:	mult	r1,r2	⇒ r1
9	e:	loadAl	r0,@y	⇒ r3
12	f:	mult	r1,r3	⇒ r1
13	g:	loadAl	r0,@z	⇒ r2
16	h:	mult	r1,r2	⇒ r1
18	i:	storeAl	r1	⇒ r0,@w
21				

The Code

⇒ **20**
cycles




The Dependence Graph

1. Build the dependence graph
2. Determine priorities: longest latency-weighted path

a:	loadAl	r0,@w	⇒ r1
b:	add	r1,r1	⇒ r1
c:	loadAl	r0,@x	⇒ r2
d:	mult	r1,r2	⇒ r1
e:	loadAl	r0,@y	⇒ r3
f:	mult	r1,r3	⇒ r1
g:	loadAl	r0,@z	⇒ r2
h:	mult	r1,r2	⇒ r1
i:	storeAl	r1	⇒ r0,@w

The Code




The Dependence Graph

1. Build the dependence graph
2. Determine priorities: longest latency-weighted path

a:	loadAl	r0,@w	⇒ r1
b:	add	r1,r1	⇒ r1
c:	loadAl	r0,@x	⇒ r2
d:	mult	r1,r2	⇒ r1
e:	loadAl	r0,@y	⇒ r3
f:	mult	r1,r3	⇒ r1
g:	loadAl	r0,@z	⇒ r2
h:	mult	r1,r2	⇒ r1
i:	storeAl	r1	⇒ r0,@w

The Code




The Dependence Graph

Note: Here we assume that operation has to finish to satisfy an anti dependence.
Our ILOC simulator takes only one cycle to satisfy an anti dependence since read-stage is executed before write stage

1. Build the dependence graph
2. Determine priorities: longest latency-weighted path

a:	loadAl	r0,@w	⇒ r1
b:	add	r1,r1	⇒ r1
c:	loadAl	r0,@x	⇒ r2
d:	mult	r1,r2	⇒ r1
e:	loadAl	r0,@y	⇒ r3
f:	mult	r1,r3	⇒ r1
g:	loadAl	r0,@z	⇒ r2
h:	mult	r1,r2	⇒ r1
i:	storeAl	r1	⇒ r0,@w

The Code




The Dependence Graph

Note: Here we assume that operation has to finish to satisfy an anti dependence.
 Our ILOC simulator **takes only one cycle to satisfy an anti dependence** since read-stage is executed before write stage (EaC).

1. Build the dependence graph
2. Determine priorities: longest latency-weighted path
3. Perform list scheduling (forward)

a:	loadAl	r0,@w	\Rightarrow r1
b:	add	r1,r1	\Rightarrow r1
c:	loadAl	r0,@x	\Rightarrow r2
d:	mult	r1,r2	\Rightarrow r1
e:	loadAl	r0,@y	\Rightarrow r3
f:	mult	r1,r3	\Rightarrow r1
g:	loadAl	r0,@z	\Rightarrow r2
h:	mult	r1,r2	\Rightarrow r1
i:	storeAl	r1	\Rightarrow r0,@w

The Code




The Dependence Graph

1. Build the dependence graph
2. Determine priorities: longest latency-weighted path
3. Perform list scheduling (forward)

1	a:	loadAl	r0,@w	⇒ r1
2	c:	loadAl	r0,@x	⇒ r2
3	e:	loadAl	r0,@y	⇒ r3
4	b:	add	r1,r1	⇒ r1
5	d:	mult	r1,r2	⇒ r1
7	g:	loadAl	r0,@z	⇒ r2
8	f:	mult	r1,r3	⇒ r1
10	h:	mult	r1,r2	⇒ r1
12	i:	storeAl	r1	⇒ r0,@w
15				

The Code

⇒ 14
cycles



The Dependence Graph

Our ILOC simulator takes only one cycle to satisfy an anti dependence

Forward list scheduling

- start with available ops
- work forward
- ready \Rightarrow all operands available

Backward list scheduling

- start with no successors
- work backward
- ready \Rightarrow latency covers operands

Different heuristics (forward) based on [Precedence/Dependence Graph](#)

1. Longest latency weighted path to root (\Rightarrow critical path)
2. Highest latency instructions (\Rightarrow more overlap)
3. Most immediate successors (\Rightarrow create more candidates)
4. Most descendents (\Rightarrow create more candidates)
5. ...

Interactions with register allocation

- perform dynamic register renaming (\Rightarrow may require spill code)
- move life ranges around (\Rightarrow may remove or require spill code)
- ...

Lexical Analysis

Read EaC: Chapters 2.1 - 2.5